

# Dealing with Failures

- Map worker failure
  - Map tasks completed or in-progress at worker are reset to idle
  - Reduce workers are notified when task is rescheduled on another worker
- Reduce worker failure
  - Only in-progress tasks are reset to idle
  - Reduce task is restarted
- Master failure
  - MapReduce task is aborted and client is notified

# How many Map and Reduce jobs?

- M map tasks, R reduce tasks
- Rule of a thumb:
  - Make M much larger than the number of nodes in the cluster
  - One DFS chunk per map is common
  - Improves dynamic load balancing and speeds up recovery from worker failures
  - Usually R is smaller than M, because output is spread across R files

# Refinements: Backup Tasks

## ■ Problem

- Slow workers significantly lengthen the job completion time:
  - Other jobs on the machine
  - Bad disks
  - Weird things

## ■ Solution

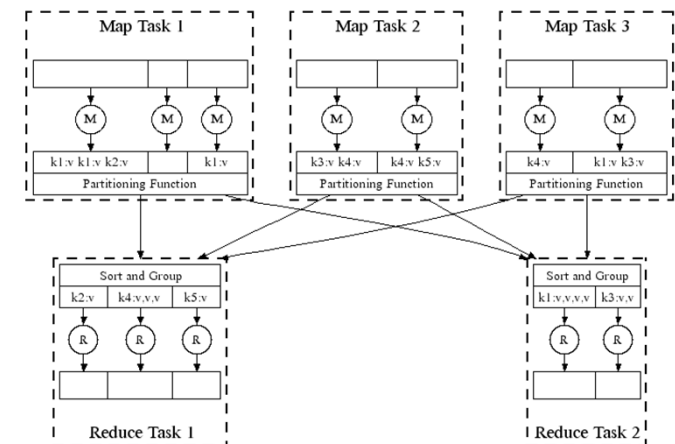
- Near end of phase, spawn backup copies of tasks
  - Whichever one finishes first “wins”

## ■ Effect

- Dramatically shortens job completion time

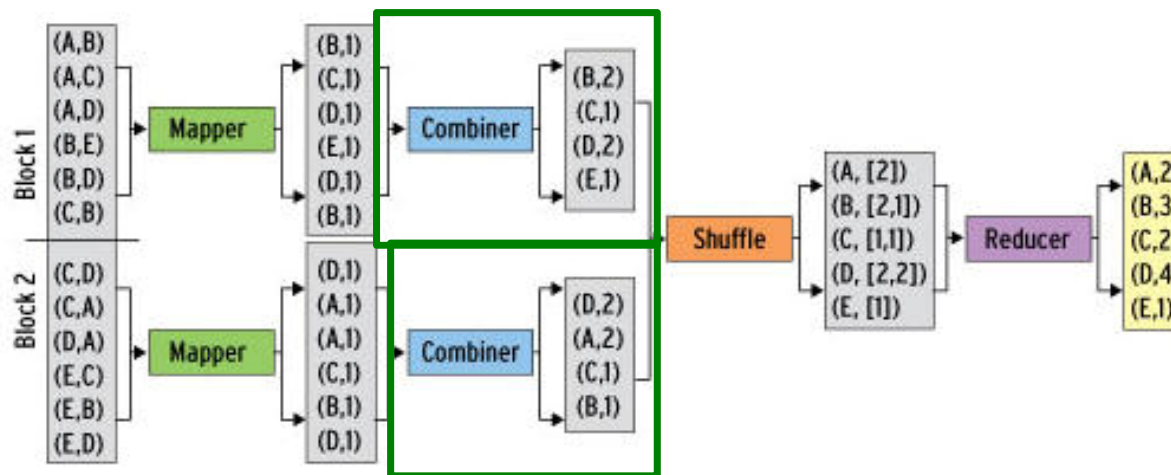
# Refinement: Combiners

- Often a Map task will produce many pairs of the form  $(k, v1)$ ,  $(k, v2)$ , ... for the same key  $k$ 
  - E.g., popular words in the word count example
- Can save network time by pre-aggregating values in the mapper:
  - $\text{combine}(k, \text{list}(v1)) \rightarrow v2$
  - Combiner is usually same as the reduce function
- Works only if reduce function is commutative and associative  
交换律和结合律, e.g., sum



# Refinement: Combiners

- Back to our word counting example:
  - Combiner combines the values of all keys of a single mapper (single machine):



- Much less data needs to be copied and shuffled!

# Refinement: Partition Function

- Want to control how keys get partitioned
  - Inputs to map tasks are created by contiguous splits of input file
  - Reduce needs to ensure that records with the same intermediate key end up at the same worker
- System uses a default partition function:
  - $\text{hash}(\text{key}) \bmod R$
- Sometimes useful to override the hash function:
  - E.g.,  $\text{hash}(\text{hostname}(\text{URL})) \bmod R$  ensures URLs from a host end up in the same output file

# Example: Host size

- Suppose we have a large web corpus (语料库) with a metadata file formatted as follows:
  - Each record of the form: (URL, size, date, ...)
- We want to: For each host(not each URL), find the total number of bytes
  - That is, the sum of the page sizes for all URLs from that particular host
- Map: For each record, output(hostname(URL),size)
- Reduce: sum the size of each host

# Example: Language Model

- Statistical machine translation:
  - Need to count number of times every 5-word sequence occurs in a large corpus of documents
- Very easy with MapReduce:
  - Map:
    - Extract (5-word sequence, count) from document
  - Reduce:
    - Combine the counts



# Example: Join By Map-Reduce

- Compute the natural join  $R(A,B) \bowtie S(B,C)$ .  $R$  and  $S$  are each stored in files. Tuples are pairs  $(a,b)$  or  $(b,c)$
- **Map:**  $(b,(R,a))$  for each tuple on  $R$ ;  $(b,(S,c))$  for each tuple on  $S$
- **Reduce:** same key with  $(R,a)$  or  $(S,c)$ , then output only  $(a,c)$ .  
key is irrelevant.

A	B
$a_1$	$b_1$
$a_2$	$b_1$
$a_3$	$b_2$
$a_4$	$b_3$

$\bowtie$

B	C
$b_2$	$c_1$
$b_2$	$c_2$
$b_3$	$c_3$

$=$

A	C
$a_3$	$c_1$
$a_3$	$c_2$
$a_4$	$c_3$

S

# Pointers and Further Reading

# Reading

- Jeffrey Dean and Sanjay Ghemawat: MapReduce: Simplified Data Processing on Large Clusters
  - <http://labs.google.com/papers/mapreduce.html>
- Sanjay Ghemawat, Howard Gobioff, and Shun-Tak Leung: The Google File System
  - <http://labs.google.com/papers/gfs.html>

# Resources

## ■ Hadoop Wiki

### ■ Introduction

- <http://wiki.apache.org/lucene-hadoop/>

### ■ Getting Started

- <http://wiki.apache.org/lucene-hadoop/GettingStartedWithHadoop>

### ■ Map/Reduce Overview

- <http://wiki.apache.org/lucene-hadoop/HadoopMapReduce>
- <http://wiki.apache.org/lucene-hadoop/HadoopMapRedClasses>

### ■ Eclipse Environment

- <http://wiki.apache.org/lucene-hadoop/EclipseEnvironment>

## ■ Javadoc

- <http://lucene.apache.org/hadoop/docs/api/>

# Resources

- Releases from Apache download mirrors
  - <http://www.apache.org/dyn/closer.cgi/lucene/hadoop/>
- Nightly builds of source
  - <http://people.apache.org/dist/lucene/hadoop/nightly/>
- Source code from subversion
  - [http://lucene.apache.org/hadoop/version\\_control.html](http://lucene.apache.org/hadoop/version_control.html)

# Further Reading

- Programming model inspired by functional language primitives
- Partitioning/shuffling similar to many large-scale sorting systems
  - NOW-Sort ['97]
- Re-execution for fault tolerance
  - BAD-FS ['04] and TACC ['97]
- Locality optimization has parallels with Active Disks/Diamond work
  - Active Disks ['01], Diamond ['04]
- Backup tasks similar to Eager Scheduling in Charlotte system
  - Charlotte ['96]
- Dynamic load balancing solves similar problem as River's distributed queues
  - River ['99]